

Lesson 3 GridSearchCV的基础及进阶使用方法

1.sklearn中GridSearchCV的使用方法

```
In [1]: # 科学计算模块
import numpy as np
import pandas as pd

# Scikit-Learn相关模块
# 评估器类
from sklearn.linear_model import LogisticRegression

# 实用函数
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
```

```
In [2]: from sklearn.model_selection import GridSearchCV
```

```
In [39]: GridSearchCV?
```

```
Init signature: GridSearchCV(estimator, param_grid, *, scoring=None, n_jobs=None,
refit=True, cv=None, verbose=0, pre_dispatch='2*n_jobs', error_score=nan, return_t
rain_score=False)
```

Docstring:

Exhaustive search over specified parameter values for an estimator.

Important members are fit, predict.

GridSearchCV implements a "fit" and a "score" method.
It also implements "score_samples", "predict", "predict_proba",
"decision_function", "transform" and "inverse_transform" if they are
implemented in the estimator used.

The parameters of the estimator used to apply these methods are optimized
by cross-validated grid-search over a parameter grid.

Read more in the :ref:`User Guide <grid_search>`.

Parameters

estimator : estimator object.

This is assumed to implement the scikit-learn estimator interface.
Either estimator needs to provide a ``score`` function,
or ``scoring`` must be passed.

param_grid : dict or list of dictionaries

Dictionary with parameters names (``str``) as keys and lists of
parameter settings to try as values, or a list of such
dictionaries, in which case the grids spanned by each dictionary
in the list are explored. This enables searching over any sequence
of parameter settings.

scoring : str, callable, list, tuple or dict, default=None

Strategy to evaluate the performance of the cross-validated model on
the test set.

If ``scoring`` represents a single score, one can use:

- a single string (see :ref:`scoring_parameter`);
- a callable (see :ref:`scoring`) that returns a single value.

If ``scoring`` represents multiple scores, one can use:

- a list or tuple of unique strings;
- a callable returning a dictionary where the keys are the metric
names and the values are the metric scores;
- a dictionary with metric names as keys and callables as values.

See :ref:`multimetric_grid_search` for an example.

n_jobs : int, default=None

Number of jobs to run in parallel.
``None`` means 1 unless in a :obj:`joblib.parallel_backend` context.
``-1`` means using all processors. See :term:`Glossary <n_jobs>`
for more details.

.. versionchanged:: v0.20

``n_jobs`` default changed from 1 to None

refit : bool, str, or callable, default=True

Refit an estimator using the best found parameters on the whole
dataset.

For multiple metric evaluation, this needs to be a ``str`` denoting the
scorer that would be used to find the best parameters for refitting

the estimator at the end.

Where there are considerations other than maximum score in choosing a best estimator, `refit` can be set to a function which returns the selected `best_index_` given `cv_results_`. In that case, the `best_estimator_` and `best_params_` will be set according to the returned `best_index_` while the `best_score_` attribute will not be available.

The refitted estimator is made available at the `best_estimator_` attribute and permits using `predict` directly on this `GridSearchCV` instance.

Also for multiple metric evaluation, the attributes `best_index_`, `best_score_` and `best_params_` will only be available if `refit` is set and all of them will be determined w.r.t this specific scorer.

See `scoring` parameter to know more about multiple metric evaluation.

.. versionchanged:: 0.20
Support for callable added.

`cv` : int, cross-validation generator or an iterable, default=None
Determines the cross-validation splitting strategy.
Possible inputs for `cv` are:

- None, to use the default 5-fold cross validation,
- integer, to specify the number of folds in a `(Stratified)KFold`,
- :term:`CV splitter`,
- An iterable yielding (train, test) splits as arrays of indices.

For integer/None inputs, if the estimator is a classifier and `y` is either binary or multiclass, `:class:`StratifiedKFold`` is used. In all other cases, `:class:`KFold`` is used. These splitters are instantiated with `shuffle=False` so the splits will be the same across calls.

Refer `:ref:`User Guide <cross_validation>`` for the various cross-validation strategies that can be used here.

.. versionchanged:: 0.22
`cv` default value if None changed from 3-fold to 5-fold.

`verbose` : int
Controls the verbosity: the higher, the more messages.

- >1 : the computation time for each fold and parameter candidate is displayed;
- >2 : the score is also displayed;
- >3 : the fold and candidate parameter indexes are also displayed together with the starting time of the computation.

`pre_dispatch` : int, or str, default=n_jobs
Controls the number of jobs that get dispatched during parallel execution. Reducing this number can be useful to avoid an explosion of memory consumption when more jobs get dispatched than CPUs can process. This parameter can be:

- None, in which case all the jobs are immediately created and spawned. Use this for lightweight and fast-running jobs, to avoid delays due to on-demand spawning of the jobs
- An int, giving the exact number of total jobs that are

spawned

- A str, giving an expression as a function of n_jobs, as in '2*n_jobs'

error_score : 'raise' or numeric, default=np.nan

Value to assign to the score if an error occurs in estimator fitting. If set to 'raise', the error is raised. If a numeric value is given, FitFailedWarning is raised. This parameter does not affect the refit step, which will always raise the error.

return_train_score : bool, default=False

If ``False``, the ``cv_results_`` attribute will not include training scores.

Computing training scores is used to get insights on how different parameter settings impact the overfitting/underfitting trade-off. However computing the scores on the training set can be computationally expensive and is not strictly required to select the parameters that yield the best generalization performance.

.. versionadded:: 0.19

.. versionchanged:: 0.21

Default value was changed from ``True`` to ``False``

Examples

```

>>> from sklearn import svm, datasets
>>> from sklearn.model_selection import GridSearchCV
>>> iris = datasets.load_iris()
>>> parameters = {'kernel':('linear', 'rbf'), 'C':[1, 10]}
>>> svc = svm.SVC()
>>> clf = GridSearchCV(svc, parameters)
>>> clf.fit(iris.data, iris.target)
GridSearchCV(estimator=SVC(),
              param_grid={'C': [1, 10], 'kernel': ('linear', 'rbf')})
>>> sorted(clf.cv_results_.keys())
['mean_fit_time', 'mean_score_time', 'mean_test_score', ...
 'param_C', 'param_kernel', 'params', ...
 'rank_test_score', 'split0_test_score', ...
 'split2_test_score', ...
 'std_fit_time', 'std_score_time', 'std_test_score']

```

Attributes

cv_results_ : dict of numpy (masked) ndarrays

A dict with keys as column headers and values as columns, that can be imported into a pandas ``DataFrame``.

For instance the below given table

param_kernel	param_gamma	param_degree	split0_test_score	...	rank_t...
'poly'	--	2	0.80	...	2
'poly'	--	3	0.70	...	4
'rbf'	0.1	--	0.80	...	3
'rbf'	0.2	--	0.93	...	1

will be represented by a ``cv_results_`` dict of::

```
{
  'param_kernel': masked_array(data = ['poly', 'poly', 'rbf', 'rbf'],
                                mask = [False False False False]...)
  'param_gamma': masked_array(data = [-- -- 0.1 0.2],
                                mask = [ True  True False False]...),
  'param_degree': masked_array(data = [2.0 3.0 -- --],
                                mask = [False False  True  True]...),
  'split0_test_score' : [0.80, 0.70, 0.80, 0.93],
  'split1_test_score' : [0.82, 0.50, 0.70, 0.78],
  'mean_test_score'    : [0.81, 0.60, 0.75, 0.85],
  'std_test_score'     : [0.01, 0.10, 0.05, 0.08],
  'rank_test_score'    : [2, 4, 3, 1],
  'split0_train_score' : [0.80, 0.92, 0.70, 0.93],
  'split1_train_score' : [0.82, 0.55, 0.70, 0.87],
  'mean_train_score'   : [0.81, 0.74, 0.70, 0.90],
  'std_train_score'    : [0.01, 0.19, 0.00, 0.03],
  'mean_fit_time'      : [0.73, 0.63, 0.43, 0.49],
  'std_fit_time'       : [0.01, 0.02, 0.01, 0.01],
  'mean_score_time'    : [0.01, 0.06, 0.04, 0.04],
  'std_score_time'     : [0.00, 0.00, 0.00, 0.01],
  'params'             : [{'kernel': 'poly', 'degree': 2}, ...],
}
```

NOTE

The key ``'params'`` is used to store a list of parameter settings dicts for all the parameter candidates.

The ``mean_fit_time``, ``std_fit_time``, ``mean_score_time`` and ``std_score_time`` are all in seconds.

For multi-metric evaluation, the scores for all the scorers are available in the ``cv_results`` dict at the keys ending with that scorer's name (``'_<scorer_name>'``) instead of ``'_score'`` shown above. ('split0_test_precision', 'mean_train_precision' etc.)

best_estimator_ : estimator

Estimator that was chosen by the search, i.e. estimator which gave highest score (or smallest loss if specified) on the left out data. Not available if ``refit=False``.

See ``refit`` parameter for more information on allowed values.

best_score_ : float

Mean cross-validated score of the best_estimator

For multi-metric evaluation, this is present only if ``refit`` is specified.

This attribute is not available if ``refit`` is a function.

best_params_ : dict

Parameter setting that gave the best results on the hold out data.

For multi-metric evaluation, this is present only if ``refit`` is specified.

best_index_ : int

The index (of the ``cv_results`` arrays) which corresponds to the best candidate parameter setting.

The dict at ``search.cv_results_['params'][search.best_index_]`` gives the parameter setting for the best model, that gives the highest mean score (``search.best_score``).

For multi-metric evaluation, this is present only if ``refit`` is specified.

scorer_ : function or a dict
Scorer function used on the held out data to choose the best parameters for the model.

For multi-metric evaluation, this attribute holds the validated ``scoring`` dict which maps the scorer key to the scorer callable.

n_splits_ : int
The number of cross-validation splits (folds/iterations).

refit_time_ : float
Seconds used for refitting the best model on the whole dataset.

This is present only if ``refit`` is not False.

.. versionadded:: 0.20

multimetric_ : bool
Whether or not the scorers compute several metrics.

Notes

The parameters selected are those that maximize the score of the left out data, unless an explicit score is passed in which case it is used instead.

If ``n_jobs`` was set to a value higher than one, the data is copied for each point in the grid (and not ``n_jobs`` times). This is done for efficiency reasons if individual jobs take very little time, but may raise errors if the dataset is large and not enough memory is available. A workaround in this case is to set ``pre_dispatch``. Then, the memory is copied only ``pre_dispatch`` many times. A reasonable value for ``pre_dispatch`` is ``2 * n_jobs``.

See Also

ParameterGrid : Generates all the combinations of a hyperparameter grid.
train_test_split : Utility function to split the data into a development set usable for fitting a GridSearchCV instance and an evaluation set for its final evaluation.
sklearn.metrics.make_scorer : Make a scorer from a performance metric or loss function.

File: c:\users\user\anaconda3\lib\site-packages\sklearn\model_selection_search.py
Type: ABCMeta

Name	Description
estimator	调参对象，某评估器
param_grid	参数空间，可以是字典或者字典构成的列表，稍后介绍参数空间的创建方法
scoring	评估指标，支持同时输出多个参数
n_jobs	设置工作时参与计算的CPU逻辑核数
refit	挑选评估指标和最佳参数，在完整数据集上进行训练
cv	交叉验证的折数
verbose	输出工作日志形式
pre_dispatch	多任务并行时任务划分数量

Name	Description
error_score	当网格搜索报错时返回结果，选择'raise'时将直接报错并中断训练过程，其他情况会显示警告信息后继续完成训练
return_train_score	在交叉验证中是否显示训练集中参数得分

- `n_jobs`：设置工作时参与计算的CPU逻辑核数。值为-1表示使用全部CPU资源进行并行计算。

CPU个数、核数、逻辑核数的概念辨析

(1) CPU个数：电脑插槽上的CPU个数, 物理cpu数量

(2) CPU核数：一个物理CPU上面能处理数据的芯片组的数量


(3) CPU逻辑核数：一般情况，我们认为一颗cpu可以有多核，加上intel的超线程技术(HT), 可以在逻辑上把一个物理线程模拟出两个线程来使用，使得单个核心用起来像两个核一样，以充分发挥CPU的性能


-> 总核数 = 物理CPU个数 X 每个物理CPU的核数

-> 总逻辑核数= 物理CPU个数 X 每个物理CPU的核数 X 超线程数

- 关于CPU逻辑核数的查看方法：

(1) 可以打开“任务管理器”-“性能”窗口中，在“cpu使用记录”中看到四个窗口，表示有四个逻辑核数。win10系统可在“性能”中看到“逻辑处理器”数量为16，则表示有16个逻辑核数。





(2) 查看CPU核数：在cmd命令中输入wmic，然后在出现的新窗口中输入cpu get NumberOfCores

查看CPU逻辑核数：在cmd命令中输入wmic，然后在出现的新窗口中输入cpu get NumberOfLogicalProcessors

```
C:\Users\lenovo>wmic
wmic:root\cli>cpu get NumberOfCores
NumberOfCores
2

wmic:root\cli>cpu get NumberOfLogicalProcessors
NumberOfLogicalProcessors
4
```

1.1 GridSearchCV评估器训练过程

• Step 1.创建评估器

首先我们还是需要实例化一个评估器，这里可以是一个模型、也可以是一个机器学习流，网格搜索都可以对其进行调参。此处我们先从简单入手，尝试实例化逻辑回归模型并对其进行调参。

```
In [40]: # 数据导入
X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=24)
```



```
In [3]: clf = LogisticRegression(max_iter=int(1e6), solver='saga')
```

此处将solver设置成saga，也是为了方便后续同时比较l1正则化和l2正则化时无需更换求解器。

```
In [42]: clf.get_params()
```

```
Out[42]: {'C': 1.0,
          'class_weight': None,
          'dual': False,
          'fit_intercept': True,
          'intercept_scaling': 1,
          'l1_ratio': None,
          'max_iter': 1000000,
          'multi_class': 'auto',
          'n_jobs': None,
          'penalty': 'l2',
          'random_state': None,
          'solver': 'saga',
          'tol': 0.0001,
          'verbose': 0,
          'warm_start': False}
```

- Step 2.创建参数空间

此处我们挑选penalty和C这两个参数来进行参数空间的构造。参数空间首先可以是一个字典：

```
In [43]: param_grid_simple = {'penalty': ['l1', 'l2'],
                              'C': [1, 0.5, 0.1, 0.05, 0.01]}
```

其中，字典的Key用参数的字符串来代表不同的参数，对应的Value则用列表来表示对应参数不同的取值范围。也就是字典的Key是参数空间的维度，而Value则是不同纬度上可选的取值。而后续的网格搜索则是在上述参数的不同组合中挑选出一组最优的参数取值。

我们可以创造多个参数空间（字典），然后将其封装在一个列表中，而该列表则表示多个参数空间的集成。

```
In [44]: param_grid_multi = [
          {'penalty': ['l1', 'l2'], 'C': [1, 0.5, 0.1, 0.05, 0.01]},
          {'penalty': ['elasticnet'], 'C': [1, 0.5, 0.1, 0.05, 0.01], 'l1_ratio': [0.3, 0.6,
          ]}
```

- Step 3.实例化网格搜索评估器

```
In [45]: search = GridSearchCV(estimator=clf,
                               param_grid=param_grid_simple)
```

- Step 4.训练网格搜索评估器

```
In [46]: search.fit(X_train, y_train)
```

```
Out[46]: GridSearchCV(estimator=LogisticRegression(max_iter=1000000, solver='saga'),
                      param_grid={'C': [1, 0.5, 0.1, 0.05, 0.01],
                                   'penalty': ['l1', 'l2']})
```

1.2 GridSearchCV评估器结果查看

Name	Description
cvresults	交叉验证过程中的重要结果
bestestimator	最终挑选出的最优
bestscore	在最优参数情况下，训练集的交叉验证的平均得分
bestparams	最优参数组合
bestindex	CV过程会对所有参数组合标号，该参数表示最优参数组合的标号
scorer	在最优参数下，计算模型得分的方法
nsplits	交叉验证的折数

```
In [47]: #best_estimator_: 训练完成后的最佳评估器
search.best_estimator_

Out[47]: LogisticRegression(C=1, max_iter=1000000, penalty='l1', solver='saga')

In [48]: #best_score_: 最优参数时交叉验证时验证集准确率的平均值，而不是所有数据的准确率
search.best_score_

Out[48]: 0.9644268774703558

In [49]: search.best_params_ #最优参数组合

Out[49]: {'C': 1, 'penalty': 'l1'}

In [50]: search.best_index_ ##最优参数组合的索引

Out[50]: 0

In [51]: # 查看参数
search.best_estimator_.coef_

Out[51]: array([[ 0.          ,  0.          , -3.4734417 ,  0.          ],
       [ 0.          ,  0.          ,  0.          ,  0.          ],
       [-0.5550736 , -0.34229207,  3.03236645,  4.12146594]])

In [52]: # 查看训练误差、测试误差
search.best_estimator_.score(X_train,y_train), search.best_estimator_.score(X_test,y_t

Out[52]: (0.9732142857142857, 0.9736842105263158)

In [53]: # 等价于search.best_estimator_.score
search.score(X_train,y_train), search.score(X_test,y_test)

Out[53]: (0.9732142857142857, 0.9736842105263158)

In [54]: search.n_splits_

Out[54]: 5

In [55]: search.refit_time_ #在整个训练集上训练最佳模型的用时

Out[55]: 0.10850024223327637

In [56]: search.cv_results_

```

```

Out[56]: {'mean_fit_time': array([0.08030992, 0.04200077, 0.05280018, 0.02789931, 0.01040025,
                                0.01521077, 0.00429974, 0.01109977, 0.00080037, 0.00449982]),
          'std_fit_time': array([8.12094464e-03, 1.24215189e-02, 2.15884069e-03, 7.34262965e-04,
                                6.63032876e-04, 3.94517573e-04, 2.44776570e-04, 1.99462722e-04,
                                2.44697944e-04, 3.56832255e-07]),
          'mean_score_time': array([3.99827957e-04, 3.99780273e-04, 3.99780273e-04, 2.00319290e-04,
                                    2.99882889e-04, 3.00693512e-04, 9.98973846e-05, 3.00407410e-04,
                                    1.99747086e-04, 1.99937820e-04]),
          'std_score_time': array([0.00019993, 0.00019989, 0.00019989, 0.00024534, 0.00024485,
                                    0.00024552, 0.00019979, 0.00024528, 0.00024464, 0.00024487]),
          'param_C': masked_array(data=[1, 1, 0.5, 0.5, 0.1, 0.1, 0.05, 0.05, 0.01, 0.01],
                                   mask=[False, False, False, False, False, False, False, False, False, False],
                                   fill_value='?',
                                   dtype=object),
          'param_penalty': masked_array(data=['l1', 'l2', 'l1', 'l2', 'l1', 'l2', 'l1', 'l2', 'l1', 'l2'],
                                         mask=[False, False, False, False, False, False, False, False, False, False],
                                         fill_value='?',
                                         dtype=object),
          'params': [{'C': 1, 'penalty': 'l1'},
                     {'C': 1, 'penalty': 'l2'},
                     {'C': 0.5, 'penalty': 'l1'},
                     {'C': 0.5, 'penalty': 'l2'},
                     {'C': 0.1, 'penalty': 'l1'},
                     {'C': 0.1, 'penalty': 'l2'},
                     {'C': 0.05, 'penalty': 'l1'},
                     {'C': 0.05, 'penalty': 'l2'},
                     {'C': 0.01, 'penalty': 'l1'},
                     {'C': 0.01, 'penalty': 'l2'}],
          'split0_test_score': array([1.          , 1.          , 1.          , 1.          , 1.
                                     , 1.          , 0.82608696, 1.          , 0.30434783, 0.91304348]),
          'split1_test_score': array([0.91304348, 0.91304348, 0.82608696, 0.86956522, 0.8260869
                                     6, 0.73913043, 0.69565217, 0.73913043, 0.39130435, 0.69565217]),
          'split2_test_score': array([1.          , 1.          , 1.          , 1.          , 0.9545454
                                     5, 0.95454545, 0.86363636, 0.90909091, 0.36363636, 0.86363636]),
          'split3_test_score': array([0.95454545, 0.95454545, 0.95454545, 0.90909091, 0.9545454
                                     5, 0.95454545, 0.86363636, 0.90909091, 0.36363636, 0.90909091]),
          'split4_test_score': array([0.95454545, 0.95454545, 0.95454545, 0.95454545, 0.9545454
                                     5, 0.90909091, 0.86363636, 0.95454545, 0.36363636, 0.90909091]),
          'mean_test_score': array([0.96442688, 0.96442688, 0.94703557, 0.94664032, 0.93794466,
                                    0.91146245, 0.82252964, 0.90237154, 0.35731225, 0.85810277]),
          'std_test_score': array([0.03276105, 0.03276105, 0.06379941, 0.05120065, 0.05863407,
                                    0.09083516, 0.06508431, 0.08830786, 0.02856808, 0.08323326]),
          'rank_test_score': array([ 1,  1,  3,  4,  5,  6,  9,  7, 10,  8])}

```

```

In [57]: pd.DataFrame(search.cv_results_) #可以pandas的DataFrame形式展示，rank_test_score为验证

```

Out[57]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_penalty	params
0	0.080310	8.120945e-03	0.000400	0.000200	1	l1	{'C': 1, 'penalty': 'l1'}
1	0.042001	1.242152e-02	0.000400	0.000200	1	l2	{'C': 1, 'penalty': 'l2'}
2	0.052800	2.158841e-03	0.000400	0.000200	0.5	l1	{'C': 0.5, 'penalty': 'l1'}
3	0.027899	7.342630e-04	0.000200	0.000245	0.5	l2	{'C': 0.5, 'penalty': 'l2'}
4	0.010400	6.630329e-04	0.000300	0.000245	0.1	l1	{'C': 0.1, 'penalty': 'l1'}
5	0.015211	3.945176e-04	0.000301	0.000246	0.1	l2	{'C': 0.1, 'penalty': 'l2'}
6	0.004300	2.447766e-04	0.000100	0.000200	0.05	l1	{'C': 0.05, 'penalty': 'l1'}
7	0.011100	1.994627e-04	0.000300	0.000245	0.05	l2	{'C': 0.05, 'penalty': 'l2'}
8	0.000800	2.446979e-04	0.000200	0.000245	0.01	l1	{'C': 0.01, 'penalty': 'l1'}
9	0.004500	3.568323e-07	0.000200	0.000245	0.01	l2	{'C': 0.01, 'penalty': 'l2'}

2.借助机器学习流的全域参数搜索方法

In [58]:

```
# 科学计算模块
import numpy as np
import pandas as pd

# 绘图模块
import matplotlib as mpl
import matplotlib.pyplot as plt

# Scikit-Learn相关模块
# 评估器类
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import roc_auc_score
# 实用函数
from sklearn.model_selection import train_test_split
```

```
# 数据准备
from sklearn.datasets import load_iris
```

```
In [59]: # 构造机器学习流
pipe = make_pipeline(PolynomialFeatures(),
                     StandardScaler(),
                     LogisticRegression(max_iter=int(1e6)))
```

```
In [60]: # 构造参数空间
param_grid = [
    {'polynomialfeatures__degree': np.arange(1, 5).tolist(), 'logisticregression__pena
    {'polynomialfeatures__degree': np.arange(1, 5).tolist(), 'logisticregression__pena
    {'polynomialfeatures__degree': np.arange(1, 5).tolist(), 'logisticregression__pena
]
```

```
In [61]: import sklearn
sorted(sklearn.metrics.SCORERS.keys())
```

```
Out[61]: ['accuracy',
          'adjusted_mutual_info_score',
          'adjusted_rand_score',
          'average_precision',
          'balanced_accuracy',
          'completeness_score',
          'explained_variance',
          'f1',
          'f1_macro',
          'f1_micro',
          'f1_samples',
          'f1_weighted',
          'fowlkes_mallows_score',
          'homogeneity_score',
          'jaccard',
          'jaccard_macro',
          'jaccard_micro',
          'jaccard_samples',
          'jaccard_weighted',
          'max_error',
          'mutual_info_score',
          'neg_brier_score',
          'neg_log_loss',
          'neg_mean_absolute_error',
          'neg_mean_absolute_percentage_error',
          'neg_mean_gamma_deviance',
          'neg_mean_poisson_deviance',
          'neg_mean_squared_error',
          'neg_mean_squared_log_error',
          'neg_median_absolute_error',
          'neg_root_mean_squared_error',
          'normalized_mutual_info_score',
          'precision',
          'precision_macro',
          'precision_micro',
          'precision_samples',
          'precision_weighted',
          'r2',
          'rand_score',
          'recall',
          'recall_macro',
          'recall_micro',
          'recall_samples',
          'recall_weighted',
          'roc_auc',
          'roc_auc_ovo',
          'roc_auc_ovo_weighted',
          'roc_auc_ovr',
          'roc_auc_ovr_weighted',
          'top_k_accuracy',
          'v_measure_score']
```

```
In [62]: search1 = GridSearchCV(estimator=pipe,
                                param_grid=param_grid,
                                n_jobs=2,
                                verbose=1,
                                scoring='roc_auc_ovr')
```

```
In [63]: search1.fit(X_train, y_train)
```

Fitting 5 folds for each of 1064 candidates, totalling 5320 fits

```

Out[63]: GridSearchCV(estimator=Pipeline(steps=[('polynomialfeatures',
                                              PolynomialFeatures()),
                                              ('standardscaler', StandardScaler()),
                                              ('logisticregression',
                                              LogisticRegression(max_iter=1000000))]),
                    n_jobs=2,
                    param_grid=[{'logisticregression__C': [0.1, 0.2,
                                                              0.30000000000000004, 0.4,
                                                              0.5, 0.6,
                                                              0.7000000000000001, 0.8,
                                                              0.9, 1.0, 1.1,
                                                              1.2000000000000002,
                                                              1.3000000000000003,
                                                              1.4000000...,
                                                              1.4000000000000001,
                                                              1.5000000000000002, 1.6,
                                                              1.7000000000000002,
                                                              1.8000000000000003,
                                                              1.9000000000000001],
                                'logisticregression__l1_ratio': [0.1, 0.2,
                                                                  0.30000000000000004,
                                                                  0.4, 0.5, 0.6,
                                                                  0.7000000000000001,
                                                                  0.8, 0.9],
                                'logisticregression__penalty': ['elasticnet'],
                                'logisticregression__solver': ['saga'],
                                'polynomialfeatures__degree': [1, 2, 3, 4]}],
                    scoring='roc_auc_ovr', verbose=True)

```

```
In [64]: search1.best_score_
```

```
Out[64]: 0.9987103174603174
```

```
In [65]: search1.best_params_
```

```
Out[65]: {'logisticregression__C': 1.5000000000000002,
          'logisticregression__penalty': 'l1',
          'logisticregression__solver': 'saga',
          'polynomialfeatures__degree': 4}
```

```
In [66]: search1.best_estimator_.score(X_train,y_train)
```

```
Out[66]: 0.9821428571428571
```

```
In [67]: search1.best_estimator_.score(X_test,y_test)
```

```
Out[67]: 0.9736842105263158
```

```

In [68]: search2 = GridSearchCV(estimator=pipe,
                                param_grid=param_grid,
                                n_jobs=2,
                                verbose=1,
                                scoring='accuracy')
search2.fit(X_train, y_train)

```

Fitting 5 folds for each of 1064 candidates, totalling 5320 fits

```

Out[68]: GridSearchCV(estimator=Pipeline(steps=[('polynomialfeatures',
                                             PolynomialFeatures()),
                                             ('standardscaler', StandardScaler()),
                                             ('logisticregression',
                                             LogisticRegression(max_iter=1000000))]),
                    n_jobs=2,
                    param_grid=[{'logisticregression__C': [0.1, 0.2,
                                                             0.30000000000000004, 0.4,
                                                             0.5, 0.6,
                                                             0.7000000000000001, 0.8,
                                                             0.9, 1.0, 1.1,
                                                             1.2000000000000002,
                                                             1.3000000000000003,
                                                             1.4000000...,
                                                             1.4000000000000001,
                                                             1.5000000000000002, 1.6,
                                                             1.7000000000000002,
                                                             1.8000000000000003,
                                                             1.9000000000000001],
                                'logisticregression__l1_ratio': [0.1, 0.2,
                                                                  0.30000000000000004,
                                                                  0.4, 0.5, 0.6,
                                                                  0.7000000000000001,
                                                                  0.8, 0.9],
                                'logisticregression__penalty': ['elasticnet'],
                                'logisticregression__solver': ['saga'],
                                'polynomialfeatures__degree': [1, 2, 3, 4]}],
                    scoring='accuracy', verbose=True)

```

```
In [69]: search2.best_score_
```

```
Out[69]: 0.9731225296442687
```

```
In [70]: search2.best_params_
```

```
Out[70]: {'logisticregression__C': 0.6,
          'logisticregression__penalty': 'l2',
          'logisticregression__solver': 'sag',
          'polynomialfeatures__degree': 3}
```

```
In [71]: search2.best_estimator_.score(X_train,y_train)
```

```
Out[71]: 0.9821428571428571
```

```
In [72]: search2.best_estimator_.score(X_test,y_test)
```

```
Out[72]: 0.9473684210526315
```